

Long-Run Dynamics of Blockwise Self-Conditioned Language Generation After Prompt Eviction

Semantic Afterlife

Anonymous

Abstract

A language model whose input is truncated to the last W generated tokens continues after the original prompt has left that window. We study the *blockwise re-prompt kernel* $Y_k^{(B)} \sim P_\theta(\cdot \mid X_k; B, \tau)$, $X_{k+1} = \text{Tail}_W(X_k \oplus Y_k^{(B)})$, with $B=1024$ and $W=4096$ ($B/W=0.25$) on one instruct generator. After full seed eviction, a textual repetition lock is the typical $T=0.3$ outcome (19/20 domain trajectories). Last-band occupancy retains between-seed distinguishability on ten fixed seed texts (F4); that gap is not recovered prompt memory (H2). Hypothesis H1 (validated Markov state models) is unsupported. We did not measure a sliding native KV cache. Degeneracy is the sample, not a filter that was applied and then forgotten.

1 Introduction

Let $X_k \in V^W$ be the last W generator tokens. The object of this paper is not token-by-token sliding attention. It is the blockwise kernel

$$Y_k^{(B)} \sim P_\theta(\cdot \mid X_k; B, \tau), \quad X_{k+1} = \text{Tail}_W(X_k \oplus Y_k^{(B)}). \quad (1)$$

After generated count $g = W$ the seed is physically absent from the model’s input (full eviction). The subsequent trajectory is the *afterlife*. The scientific question is whether that trajectory retains seed identity, occupies a domain-tagged lock, or diffuses.

That question is easy to answer wrongly. A looping trajectory occupies one point in representation space and reports a confined mean-squared displacement for reasons that have nothing to do with semantics. Degeneracy is labelled before geometry is read.

The experiments are Stages 0–6 of a prospectively specified, version-controlled sequential study with documented adaptive replanning between stages (`docs/research-plan.md`). They do not establish H1, H5, or a result for language models as a class.

Section 2 places the afterlife against verified related work. Section 3 defines the kernel, the two eviction times, and forced continuation. Section 4 reports the low- T lock. Section 5 reports last-band domain occupancy. Section 6 states what the design cannot establish. Section 7 lists remaining threats to validity.

2 Related work

We cite only sources marked VERIFIED in the project literature log.

Zekri et al. [2024] study autoregressive generation as a Markov chain on the context and, under contraction of the next-token kernel in total variation, prove existence of a unique stationary

distribution. That theorem is the right null for a unique long-run law. It is not a measurement of last-band occupancy after eviction.

Telephone-game and paraphrase protocols pass the *whole* previous text under an instruction [Perez et al., 2025, Mohamed et al., 2025, Wang et al., 2025, Geng et al., 2026]. Our state is a finite tail with explicit prompt eviction and many turnovers of an imposed W . Ko and Geiping [2026] describe interlocutor-conditioned attractors; this protocol has no interlocutor. Chen et al. [2026] separate horizon, context, and memory; we impose W by truncation and begin after the seed has left that window. Wang et al. [2026] pin the prompt and randomly evict generated reasoning tokens for KV-cache compression; we evict the seed completely and continue uninstructed.

Surface-form degeneration under maximization decoding is documented [Holtzman et al., 2020]. Self-reinforcement of sentence repetition is the closest alternative explanation of the lock [Xu et al., 2022]. Training-time collapse on recursively generated data [Shumailov et al., 2024] is a different object: we never retrain. Markov-state methodology from molecular dynamics [Wu and Noé, 2020, Paul et al., 2019] is the toolkit Stage 3 applied; the toolkit is not the result.

3 Protocol

3.1 Generator, window, and P1

Occupancy uses `or-qwen3-8b` (Qwen3-8B-Instruct via OpenRouter; served provider `Alibaba` on every completed occupancy step), $W=4096$, $T=0.3$, block $B=1024$ so $B/W=0.25$, analysis chunks of 1024 generator tokens, non-overlapping, and twelve turnovers of W . The continuation mechanism is `raw_completion`, not `assistant_prefill` and not `chat_instructed`. Protocol P1 is the re-prompt: the last W tokens are re-encoded and sent as a new prompt. P1 and the continuation mechanism are distinct axes (glossary). The input construction is not a sliding native KV cache (protocol P2).

Eviction has two times. *Eviction start* is $W - L_0$: the window first fills and the next generated token can drop a seed token. *Full eviction* is $g = W$: no seed token remains. Last-band occupancy uses turnover ≥ 1 , which is full eviction, not the fill instant. Historical generate JSONL stored `horizon_tokens = W - L_0`; analysis does not read that field for F4.

The process is *externally sustained self-conditioning after repeated termination attempts*: stop tokens are recorded and generation is continued. Survival analysis of EOS as absorbing is parked (Paper B).

3.2 Cost law

Hosted spend follows the recorded ledger. Stage 0 audit cost \$0.0128. Stage 1 generation billed \$6.32 on `s1-pilot-core-20260830T134147Z-f1cce9ec`. Stage 4 hosted generation billed \$3.4383. Stage 5 generation billed \$1.3385. Stage 6 used cached Gemini embeddings at hosted \$0. Project hosted spend at Stage 6 close was \$16.34 of a \$200 ceiling. This correctness pass spent \$0.

3.3 Degeneracy is the sample

A trajectory is labelled a textual repetition lock when calibrated repetition statistics cross a threshold derived from a human-prose reference corpus (Carroll and Darwin): mean 3-gram Jaccard 0.083 and late-window Jaccard 0.0122. The detector reports a near-recurrent textual

Table 1: Textual repetition-lock rates, Stage 2 model axis, split by temperature, $W=4096$, $n=4$ per cell, 95% Clopper–Pearson CI. Source: `artifacts/stage-2/model-axis/rates/fixed_point_rates_by_temperature.csv`. Pooled $n=8$ over $T \in \{0.3, 1.0\}$ remains in `fixed_point_rates.csv` and is not a $T=0.3$ table. Gemma 0/4 at both temperatures is the existence proof that the Stage 1 freeze is not a property of instruct models as a class.

Generator	$T=0.3$	95% CI	$T=1.0$	95% CI
or-qwen3-8b	3/4	[0.194, 0.994]	4/4	[0.398, 1]
or-gemma-4-31b	0/4	[0, 0.602]	0/4	[0, 0.602]
or-gpt-oss-120b	4/4	[0.398, 1]	4/4	[0.398, 1]
or-muse-glimmer-30b	2/4	[0.068, 0.932]	0/4	[0, 0.602]
or-gpt-oss-20b	did not complete design T			

Table 2: Looping rate by temperature and window, Stage 4, or-qwen3-8b, $n=4$ per cell, Clopper–Pearson CI. Source: `artifacts/stage-4/grid/looping_rate_by_cell.csv`.

T	$W=4096$	CI	$W=8192$	CI
≤ 1.0	4/4	[0.398, 1]	4/4	[0.398, 1]
1.5	0/4	[0, 0.602]	2/4	[0.068, 0.932]

regime (fixed point or short cycle), not an exact period-1 test. Those two numbers are not returned.

On the occupancy panel, 19/20 domain trajectories and 1/2 love trajectories meet the lock rule; 10/10 domain seeds have at least one lock. Clopper–Pearson 19/20 is [0.751, 0.999]. Occupancy is a statement about locked trajectories. Geometry on a looping row is not interpreted as a confined semantic basin.

4 Low-temperature repetition lock

Stage 1 generated 48 trajectories on or-qwen3-8b at $W=4096$, target $T=131072$ tokens (32 turnovers). The lock rule fired on 45/48 trajectories (94%), including 24/24 at $T=0.3$ and 21/24 at $T=1.0$. Figure 1 shows a domain-gap curve on that degenerate-majority panel; it is not a non-looping subsample.

The lock is not universal across hosted generators. Table 1 splits Stage 2 rates by temperature, $W=4096$, $n=4$ per cell, Clopper–Pearson. The previously published $n=8$ rows pool $T \in \{0.3, 1.0\}$: or-gemma-4-31b 0/8 ([0, 0.369]), or-gpt-oss-120b 8/8 ([0.631, 1]), or-qwen3-8b 7/8 ([0.473, 0.997]). At $T=0.3$: gemma 0/4 ([0, 0.602]), 120B 4/4 ([0.398, 1]), Qwen 3/4 ([0.194, 0.994]). At $T=1.0$: gemma 0/4, 120B 4/4, Qwen 4/4. Direction (F2) is gemma versus 120B at both temperatures. Qwen’s $T=0.3$ interval includes 0.5.

A prefill-versus-raw arm on or-qwen3-8b is 7/8 versus 8/8; Newcombe difference CI includes 0 (Fisher $p=1$).

Stage 4 crossed temperature and window with $n=4$ per cell. At $T \leq 1.0$, both $W=4096$ and $W=8192$ are 4/4 ([0.398, 1]). At $T=1.5$, $W=4096$ is 0/4 ([0, 0.602]); $W=8192$ is 2/4 ([0.068, 0.932]). That 4/4 versus 0/4 cell contrast at $W=4096$ is not a temperature law. Hypothesis H5 (a temperature window of non-looping confined motion) is absent: clean MSD α is defined only at $T=1.5$ and is subdiffusive.

Does seed identity survive the context horizon? — bge-m3

W = 4,096 tokens. Post-horizon mean gap 0.1569, trend -0.00079 per turnover. Separated at the last observed band.




Figure 1: Stage 1 domain-gap versus turnover in bge-m3 on the 48-trajectory or-qwen3-8b pilot. Degenerate-majority panel. Data: artifacts/stage-1/separation-bge-m3/seed_separation.data.parquet.

5 Last-band occupancy

Given the lock, Stage 5–6 ask whether last-band centroids of ten fixed seed texts remain distinguishable (between-seed vs within-seed). Twenty-eight trajectories at $W=4096$, $T=0.3$, $B=1024$, twelve turnovers. Generate s5-lock-occupancy-20260905T164327Z-6780902f; embed s5-embed-lock-occupancy-20260906T030125Z-eab6e484; degeneracy s5-degeneracy-20260906T030145Z-6780902f; gemini s6-embed-third-space-20260906T082301Z-588eff8f.

Let D_{within} be the mean cosine distance among last-band centroids that share a seed text, and D_{between} the mean among different seed texts. The gap is $\Delta_{\text{dom}} = D_{\text{between}} - D_{\text{within}}$. Domain here means one labelled seed instance, not a sampled population. Seeds are labelled separated when the trajectory-bootstrap 95% CI excludes 0 from above. That estimator already resamples with multiplicity (ADR-0019). An independent seed-pair matrix recompute agrees on the point to < 0.002 in all three spaces. Those published CIs are not re-derived from embeddings: named occupancy run directories are not recoverable (ADR-0021). Regenerating trajectories would mint new run_ids and new draws, not verify the published intervals.

Table 3 reports last-band F4. All three spaces are separated under that CI rule: bge-m3 0.201 [0.065, 0.332]; qwen3-embed-8b 0.390 [0.151, 0.593]; gemini-embed-001 0.150 [0.029, 0.266]. Sign agreement is the claim. Levels are not interchangeable. Gemini’s lower bound 0.029 is an NHST pass, not architecture-independence. Physics seed 1 has 47 versus 48 chunks, so its last-band D_{within} diagonal is undefined; F4 uses $n_{\text{within}}=9$ pairs. Leave-one-seed-out on the matrix keeps the gap positive in every space (minimum 0.128). Within-pair randomisation of the 54 unique finite pairs yields $p=0.0001$ in each space (0/9999 permutations as large as the observed gap, seed 0). That is a test on ten fixed seed texts, not a domain population. The last-band gap is between-seed distinguishability, not recovered prompt memory (H2) and not domain memory.

A protocol diagnostic, not a headline: domain trajectories spend a mean fill fraction 0.853, 95% CI [0.756, 0.936], in the last quarter, with mean stop-token fraction 0.359.

Looping / repetition-lock rate on or-qwen3-8b (P1 raw), $n = 4$ per cell




Figure 2: Looping rate versus temperature at two windows, 95% Clopper-Pearson CI, $n=4$. The $T=1.5$, $W=4096$ cell is the existence of a non-looping regime on this model, not a mixing time.

MSD α on non-degenerate trajectories only; $\times = n_{\text{clean}} < 2$ (undefined)




Figure 3: Clean MSD α where defined (non-looping cells only). Missing low- T points are the H5 absence.

6 What this does not show

H1 is unsupported. Stage 3 applied VAMP/PCCA+ on embeddings of the Stage 1 panel. `validated=0` in every cell of `artifacts/stage-3/dynamics/k_stability.md`. Implied n_{macro} is unstable across K and across embedding spaces. Hypothesis H1 (metastable semantic states as an MSM object) is unsupported. A short-cycle lock is not a validated macrostate.

H5 is absent. There is no low- T clean α on the Stage 4 grid.

F6 is not a headline. Matched replica pairs are counterfactual narratives, not one-token factual flips. $n=2$ per family. Published twin CI columns used a bootstrap that dropped resample multiplicities (ADR-0019) and are not reported here as intervals. Last-band *point* Δ in `gemini-embed-001` is -0.012 (all pairs) and 0.049 in `bge-m3` waterloo. That is no detected divergence / underpowered, not occupancy of one lock and not semantic collapse. An equivalence design is not claimed.

No universal attractors. One instruct generator, P1, `raw_completion`, Alibaba, $T=0.3$. Gemma's 0/8 already falsifies universality.

Table 3: F4 last-band between-seed vs within-seed gap, ten fixed seed texts, $n=20$, $n_{\text{within}}=9$. Source: `artifacts/stage-6/occupancy/domain_separation_last_band.csv`. Separated means the published 95% CI excludes 0 from above. Published 95% CIs are original analysis outputs in that CSV; source occupancy vectors are not recoverable (ADR-0021).

Space	gap	95% CI	separated
bge-m3	0.201	[0.065, 0.332]	yes
qwen3-embed-8b	0.390	[0.151, 0.593]	yes
gemini-embed-001	0.150	[0.029, 0.266]	yes

Ten-domain seed gap vs turnover on or-qwen3-8b (P1 raw), $W=4096$ $T=0.3$




Figure 4: Between-seed last-band gap versus turnover in **bge-m3** and **qwen3-embed-8b** on the Stage 5 locked panel. Last-band CIs exclude 0 from above in both spaces.

7 Limitations

P1 versus sliding attention. Every afterlife token is produced by re-prompting the last W tokens. A native sliding KV cache (P2) was not measured. No sentence here is evidence about sliding attention.

One generator, one protocol, one temperature for occupancy. F4 is **or-qwen3-8b**, P1, **raw_completion**, served provider **Alibaba**, $T=0.3$, $W=4096$, $B=1024$. P1 is the re-prompt axis; **raw_completion** is the continuation mechanism. They are not synonyms.

Instruction-tuning and register. The occupancy generator is instruct-tuned. Stage 2 already measured a reviewer register on this model. A persona/register alternative — last-band centroids cluster by residual instruction style rather than seed domain — is not excluded by F4. No hosted base model was available (Stage 0). A local **gemma-3-1b-pt** probe at $W=256$ is a different window.

Forced continuation. The object is externally sustained self-conditioning after repeated termination attempts. Last-quarter stop fraction 0.359 is a confound to name, not an EOS-absorbing measurement.

Twins, $n=2$. F6 is underpowered. Twins are not one-fact pairs. CIs from the broken bootstrap are not used as evidence.

Ten-domain last-band gap in three embedding spaces




Figure 5: F4 between-seed last-band gap versus turnover in three embedding spaces on the same 28 trajectories. Last-band CIs exclude 0; gemini lower bound is 0.029. Levels 0.201 / 0.390 / 0.150 are not interchangeable.

Degeneracy null. The lock bar is Carroll+Darwin prose. Wikipedia-scale calibration is parked.

Representation. Absolute gaps 0.201, 0.390, 0.150 disagree. Gemini lower CI 0.029 is not architecture-independence.

Provider non-determinism. Stage 2 exact-match rates were 100/60/20/20 percent across four hosted generators. Replica pairs are not guaranteed independent draws.

Occupancy vectors. Named occupancy generate and embed run directories are not recoverable (ADR-0021). Published trajectory-bootstrap CIs are the original analysis outputs in the committed CSV, not re-derived from embeddings. Independently reproducible F4 inference in this pass is the seed-pair matrix, leave-one-seed-out, and within-pair randomisation. A new generate would be a new panel, not a restore of those intervals.

Finite horizon. Occupancy uses twelve turnovers, not an infinite afterlife.

8 Discussion

What is established, on this generator and protocol: after the seed leaves W , the typical $T=0.3$ trajectory is a textual repetition lock; last-band centroids still show a between-seed last-band gap across ten fixed seed texts in three embedding spaces (not recovered prompt memory, H2).

What is not established: H1; H5; we cannot claim architecture-independence of occupancy; a result for language models as a class; P1 versus P2; mixing times; replica collapse; one-fact twin sensitivity.




Figure 6: Last-band cosine distances in **gemini-embed-001** (illustration of the pairwise matrix, not a cluster count, not a 2-D projection).

Zekri et al. [2024] remain the right uniqueness null. Seed-tagged last-band gaps on a locked panel are compatible with more than one long-run occupancy region in embedding space. That compatibility is not a proof that the token-level chain has several stationary measures.

9 Reproducibility

Experiments through Stage 6 closed on git commit **ad1f244**. Stage 7 wrote the first manuscript. This correctness pass is ADR-0019, ADR-0020, and ADR-0021 (anonymous). it does not mint **runs/s7**. Hosted spend for the pass is \$0. The GitHub tag **state-latest** is a mutable restore pointer, not a citable archival object.

Primary **run_ids**:

- Stage 1 generate: **s1-pilot-core-20260830T134147Z-f1cce9ec**
- Stage 1 degeneracy: **s1-degeneracy-20260831T070038Z-b2f6f0c9**
- Stage 1 embed: **s1-embed-pilot-core-20260831T064430Z-f05e6a68**
- Stage 1 separation: **s1-separation-bge-m3-20260831T074042Z-9bf54fbf**
- Stage 2 model axis: **s2-model-axis-20260901T015457Z-ab59afc8**
- Stage 2 mechanism: **s2-mechanism-20260901T071519Z-dfbb173a**
- Stage 2 rates: **s2-rates-20260901T125506Z-41d2e88e**
- Stage 3 dynamics: **s3-dynamics-20260901T204521Z-f21d5908**

- Stage 4: s4-w4096-new-temps-20260904T103121Z-589c8eb1, s4-w8192-20260904T120057Z-ce82ce55
- Stage 5 generate: s5-lock-occupancy-20260905T164327Z-6780902f
- Stage 6 gemini: s6-embed-third-space-20260906T082301Z-588eff8f

References

- Mingguang Chen, Licheng Wang, and Bo Qu. The horizon gap: Planning, memory, execution, training, and evaluation for long-horizon LLM agents. *arXiv preprint arXiv:2608.06663*, 2026.
- Mingmeng Geng, Amr Mohamed, Guokan Shang, Michalis Vazirgiannis, and Thierry Poibeau. Markovian generation chains in large language models. *arXiv preprint arXiv:2603.11228*, 2026.
- Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. In *International Conference on Learning Representations*, 2020.
- Ting-Wen Ko and Jonas Geiping. Attractor states emerge in multi-turn LLM conversations. *arXiv preprint arXiv:2606.30571*, 2026.
- Amr Mohamed, Mingmeng Geng, Michalis Vazirgiannis, and Guokan Shang. LLM as a broken telephone: Iterative generation distorts information. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7493–7509, Vienna, Austria, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.371. URL <https://aclanthology.org/2025.acl-long.371/>.
- Fabian Paul, Hao Wu, Maximilian Vossel, Bert L. de Groot, and Frank Noé. Identification of kinetic order parameters for non-equilibrium dynamics. *The Journal of Chemical Physics*, 150(16):164120, 2019. doi: 10.1063/1.5083627. arXiv:1811.12551.
- Jérémy Perez, Grgur Kovač, Corentin Léger, Cédric Colas, Gaia Molinaro, Maxime Derex, Pierre-Yves Oudeyer, and Clément Moulin-Frier. When LLMs play the telephone game: Cultural attractors as conceptual tools to evaluate LLMs in multi-turn settings. In *International Conference on Learning Representations*, 2025. OpenReview hash dbdea7859f1d2fc10f2c9e79b8f5ae54.
- Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024. doi: 10.1038/s41586-024-07566-y.
- Heng Wang, Jieli Qiu, Wenting Zhao, Cheng Qian, Liangwei Yang, Jiawei Han, Heng Ji, Silvio Savarese, Shelby Heinecke, and Huan Wang. Random attention: Rethinking KV cache eviction for efficient reasoning. *arXiv preprint arXiv:2609.03430*, 2026.
- Zhilin Wang, Yafu Li, Jianhao Yan, Yu Cheng, and Yue Zhang. Unveiling attractor cycles in large language models: A dynamical systems view of successive paraphrasing. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 12740–12755, Vienna, Austria, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.acl-long.624. URL <https://aclanthology.org/2025.acl-long.624/>.
- Hao Wu and Frank Noé. Variational approach for learning Markov processes from time series data. *Journal of Nonlinear Science*, 30:23–66, 2020. doi: 10.1007/s00332-019-09567-y. arXiv:1707.04659.

Jin Xu, Xiaojiang Liu, Jianhao Yan, Deng Cai, Huayang Li, and Jian Li. Learning to break the loop: Analyzing and mitigating repetitions for neural text generation. In *Advances in Neural Information Processing Systems*, 2022.

Oussama Zekri, Ambroise Odonnat, Abdelhakim Benechehab, Linus Bleistein, Nicolas Boullé, and Ievgen Redko. Large language models as Markov chains. *arXiv preprint arXiv:2410.02724*, 2024.